

A large, two-story, light-colored building with a red-tiled roof and a central tower, surrounded by green grass and trees under a clear blue sky.

# MAHARISHI UNIVERSITY of MANAGEMENT

*Engaging the Managing Intelligence of Nature*

## Computer Science Department

**CS401 Modern Programming  
Practices (MPP)  
Professor Paul Corazza**



© 2015 Maharishi University of Management, Fairfield, Iowa

All rights reserved. No part of this slide presentation may be reproduced or utilized in any form or by any means, electronic or mechanical, including photocopying or recording, or by any information storage and retrieval system, without permission in writing from Maharishi University of Management.

# Lecture 10: Best Programming Practices with Java 8

*Living Life in Accord with Natural Law*



# Wholeness Statement

Best practices in the world of OO programming are a way of ensuring quality in code. Code that adheres to best practices tends to be easier to understand, easier to maintain, more capable of adapting in the face of changing requirements and new feature requests, and more reusable for other projects. This simple theme is reflected in individual life. There are laws governing life, both physical laws and laws pertaining to all kinds of relationships and interactions. When life flows in accordance with the laws of nature, life is supported for success and fulfillment. When awareness becomes established at its deepest level, actions and behavior springing from this profound quality of awareness spontaneously are in accord with the laws of nature.

# Overview

1. **Unit-testing Stream Pipelines**
2. Introduction to Annotations
3. Using Java 8's try-with-resources
4. Handling Exceptions Arising in Stream Pipelines
5. Concurrent Processing and Parallel Streams



# How to Unit Test Stream Pipelines?

- A. The Problem. Stream pipelines, with lambdas mixed in, form a single unit; it is not obvious how to effectively unit test the pieces of the pipeline.
- B. Two Guidelines:
  - a. If the pipeline is simple enough, we can name it as an expression and unit test it directly.
  - b. If the pipeline is more complex, pieces of pipeline can be called from support methods, and the support methods can be unit-tested.

# Unit-Testing Stream Pipelines:

## Simple Expressions

- Example: Perform a unit test on the following pipeline:

```
Function<List<String>, List<String>> allToUpperCase =  
    words -> words.stream()  
        .map(word -> word.toUpperCase())  
        .collect(Collectors.toList());
```

- Can do ordinary unit testing of allToUpperCase.

```
@Test  
public void multipleWordsToUppercase() {  
    List<String> input = Arrays.asList("a", "b", "hello");  
    List<String> result = Testing.allToUpperCase.apply(input);  
    assertEquals(Arrays.asList("A", "B", "HELLO"), result);  
}
```



# Unit-Testing Stream Pipelines:

## ComplexExpressions

- **Example:** Sorts Employees first by name (ascending), then by salary in descending order

```
//difficult to test
static Function<List<Employee>, List<Employee>> employeeSorter =
    list -> list.stream()
        .sorted(Comparator.comparing((Employee e) -> e.getName())
            .thenComparing((Employee e) -> e.getSalary(),
                Comparator.reverseOrder()))
        .collect(Collectors.toList());
```

- The key point to test is whether any two Employees are being compared in the correct way.
- This can be done by replacing the lambda expression that defines the Comparator with a method reference, together with an auxiliary method that can be placed in a companion class `LibraryCompanion`.
- Demo: `lesson10.lecture.libcompanion`



**Strategy:** Refactor the pipeline using auxiliary methods that can be tested separately.

```
static Function<List<Employee>, List<Employee>> employeeSorter =  
list -> list.stream()  
    .sorted(LibraryCompanion::compareEmps)  
    .collect(Collectors.toList());
```

**//auxiliary method, used in method reference, in the class LibraryCompanion**

```
public class LibraryCompanion {  
    static Comparator<Employee> empComp  
        = Comparator.comparing((Employee e) -> e.getName())  
            .thenComparing((Employee e) -> e.getSalary(),  
                Comparator.reverseOrder());  
    public static int compareEmps(Employee e1, Employee e2) {  
        return empComp.compare(e1, e2);  
    }  
}
```

Now the key element of the original lambda can be tested directly:

```
//verify that jim comes before joe  
assert(LibraryCompanion.compareEmps(jim, joe1) < 0);  
  
//verify that joe2 comes before joe1  
assert(LibraryCompanion.compareEmps(joe2, joe1) < 0);
```

# Unit-Testing Stream Pipelines:

## ComplexExpressions

The example suggests a best practice for unit testing when lambdas are involved:

1. Replace a lambda that needs to be tested with a method reference plus an auxiliary method
2. Then you can test the auxiliary method



# Exercise 10.1

You are developing unit tests and need to decide how to test the lambda expression shown below. Should you treat it as a simple expression or a complex expression? Create a unit test to test it. All necessary files are in package `lesson10.exercise_1` in the `InClassExercises` project. Use the class `TestLambda` for your test.

```
// A list of Customers whose checking account balance is > 50,  
//sorted by customer's last name  
public static List<Customer> specialAccounts(List<Account> accounts) {  
    return accounts.stream().filter(a -> a.getBalance() > 50)  
        .map((Account a) -> a.getCustomer())  
        .sorted(Comparator.comparing((Customer c) -> c.getLastName()))  
        .collect(Collectors.toList());  
}
```

# Overview

1. Unit-testing Stream Pipelines
2. **Introduction to Annotations**
3. Using Java 8's try-with-resources
4. Handling Exceptions Arising in Stream Pipelines
5. Concurrent Processing and Parallel Streams



# What Are Annotations?

We have seen them in various contexts already:

- a. `@Test` - JUnit 4
- b. `@Override` – to indicate that a method is being overridden
- c. `@FunctionalInterface` – to indicate that an interface is functional and may be used with lambdas
- d. `@Deprecated` – to discourage use of a method or class
- e. `@SuppressWarnings` – to hide warning messages of various kinds
- f. Javadoc annotations:
  - i. `@author`
  - ii. `@since`
  - iii. `@version`

# What Are Annotations? (cont.)

2. Annotations are tags that are inserted into source code so that some tool can process them. The tools can operate on the source level, or they can process class files into which the compiler has placed annotations.
3. To benefit from annotations that you create, you need to select a *processing tool*. You need to use annotations that your processing tool understands, then apply the processing tool to your code.
  - JUnit processes its @Test annotation
  - Java compiler processes the others shown
4. Annotations can be applied to a class, a method, a variable – in fact, anywhere qualifiers like public and static may be used



# What Are Annotations? (cont.)

- Annotations may have zero or more *elements*. Here is an example of a user-defined annotation that has two elements, `assignedTo` and `severity`.

```
@BugReport(assignedTo="Harry", severity=10)  
[This annotation could be applied at the class level. Like any user-defined  
annotation, it would require an external tool to process it.]
```

- When an annotation has just one element and its name is “`value`”, the following more compact form can be used:

```
@SuppressWarnings("unchecked")  
[same as @SuppressWarnings(value = "unchecked")]
```

- If the annotations have the same type, then this is called a repeating annotation:

```
@Author(name = "Jane Doe")  
@Author(name = "John Smith")  
class MyClass { ... }
```

Repeating annotations are supported as of the Java SE 8 release

# What Are Annotations? (cont.)

8. *User-defined annotations.* The `@interface` keyword is the way the Java compiler knows you are creating an annotation; such “classes” extend the `Annotation` interface. See the demo `lesson10.lecture.annotation`

```
@Retention(RetentionPolicy.RUNTIME)
@Target(ElementType.TYPE)
public @interface BugReport {
    String assignedTo() default "[none]";
    int severity() default 0;
}
```

**Note:** the code shows how the elements "assigned to" and "severity" are defined.



# Annotation Reference

**@Retention** [@Retention](#) annotation specifies how the marked annotation is stored:

- RetentionPolicy.SOURCE – The marked annotation is retained only in the source level and is ignored by the compiler.
- RetentionPolicy.CLASS – The marked annotation is retained by the compiler at compile time, but is ignored by the Java Virtual Machine (JVM).
- RetentionPolicy.RUNTIME – The marked annotation is retained by the JVM so it can be used by the runtime environment.

**@Target** [@Target](#) annotation marks another annotation to restrict what kind of Java elements the annotation can be applied to. A target annotation specifies one of the following element types as its value:

- ElementType.ANNOTATION\_TYPE can be applied to an annotation type.
- ElementType.CONSTRUCTOR can be applied to a constructor.
- ElementType.FIELD can be applied to a field or property.
- ElementType.LOCAL\_VARIABLE can be applied to a local variable.
- ElementType.METHOD can be applied to a method-level annotation.
- ElementType.PACKAGE can be applied to a package declaration.
- ElementType.PARAMETER can be applied to the parameters of a method.
- ElementType.TYPE can be applied to any element of a class.

## Exercise 10.2

Work with the Buganizer code shown in class (available in package `lesson10.exercise_2` in `InClassExercises`). A class `StillMoreBadCode` has been added. Add to this class annotation values for the elements "assigned to" and "severity". Then integrate this new class into the `AnnotationInfo` class and run it. Info about `StillMoreBadCode` should be included in the output report.



# Overview

1. Unit-testing Stream Pipelines
2. Introduction to Annotations
3. **Using Java 8's try-with-resources**
4. Handling Exceptions Arising in Stream Pipelines
5. Concurrent Processing and Parallel Streams

# Handling Exceptions with Try-with-Resources

1. *You Wish to Use Resources Carefully.* When your application uses external resources, like files or a database connection, it is important to close the connections after your application has finished using them. Typically, using these connections involves checked exceptions; but whether or not an exception is thrown, your application must disconnect from the resource, or there can be a memory leak, causing memory to fill up.
2. *Do clean-up in a finally block.* The usual way to clean up resources is in a finally block, which will execute whether or not an exception is thrown.
3. *Complication.* But, what if closing a resource is also capable of throwing a checked exception? How should this be handled?

Demo: `lesson10.lecture.trickycatch1`



# The Closing Resources Problem

**Example:** Consider this code (lesson10.lecture.trickycatch1):

```
public void handleFile(File f) throws IOException {
    FileReader fileReader = null;
    BufferedReader buf = null;
    try{
        fileReader = new FileReader(f);
        buf = new BufferedReader(fileReader);
        String line = buf.readLine();
        System.out.println("Line from file: " + line);
    } catch(IOException e) {
        Log.config("Caught 1st IOException: " + e.getMessage());
        throw(e);
    } finally {
        try {
            if(buf != null) buf.close();
            if(fileReader != null) fileReader.close();
            throw new IOException("Caught 2nd IOException: error closing readers");
        } catch(IOException e) {
            Log.fine(e.getMessage());
            throw(e);
        }
    }
}
```

# The Closing Resources Problem (continued)

## Some Problems with the Code.

1. **First problem.** It's messy and it's easy to implement it incorrectly.
2. **More serious problem.** Consider this scenario:
  - After `BufferedReader` object is created, a call to `readLine` causes an `IOException` to be thrown.
  - In the `finally` clause, an attempt is made to close the `BufferedReader`, but the `close()` operation also throws an `IOException`.

Only one exception can be thrown, and the way Java does it is to throw the close operation exception (the one that we don't care about).

In the demo code, the message displayed will be “Caught 2<sup>nd</sup> `IOException`: error closing reader.” The `IOException` indicating a failure to execute `readLine()` is lost.

Demo: `lesson10.lecture.trickycatch1.MyClass2`



# The Try-with-Resources Solution

Here is the Java 8 alternative to the previously displayed code:

```
public void handleFile(File f) {  
    //try with resources construct  
    try (BufferedReader buf = new BufferedReader(new FileReader(f))) {  
        String line = buf.readLine();  
        System.out.println(line);  
    } catch (IOException e) {  
        log.warning("Main exception: " + e.getMessage());  
        List<Throwable> suppressed = Arrays.asList(e.getSuppressed());  
        suppressed.forEach(except -> log.warning("Suppressed message: "  
                                                + except.getMessage()));  
    }  
}
```

Demo: `lesson10.lecture.trickycatch4_trywithres.MyClass`

# The try-with-resources Solution (cont.)

## Notes

1. The resources named in the "try" clause – in this case, a `BufferedReader` – will be closed at the completion of the try block.
2. If an exception is thrown during execution of try and an error occurs in closing these resources, the close exceptions will automatically be appended to the main exception as *suppressed exceptions*.
3. If no main exception occurs, but a close exception occurs, the close exception is thrown in the usual way.
4. It is now possible to read the suppressed exceptions using the `getSuppressed()` method, which returns a list of `Throwables`. These can be read and handled if desired, but the main exception is the one that is thrown and caught.



# Exercise 10.3

Use the try-with-resources construct to rewrite the class `OldFileIO`, which reads a file `word_test.txt` using pre-Java 8 approach. The necessary files are in the `InClassExercises` project, in the `lesson10.exercise_3` package. Do the following:

1. Include a catch clause to catch the primary `IOException` if there is one.
2. In your catch clause, loop through and print any suppressed exceptions that may be thrown (as in the slides)
3. Notice how the file `word_test.txt` is accessed in `OldFileIO` (see `FILE_LOCATION` constant). Should this be improved? How?

# More About Resources

A *resource* is an object that must be closed after the program is finished with it. The `try-with-resources` statement ensures that each resource is closed at the end of the statement.

Another example of a resource that can be managed with `try-with-resources` is the `Connection` object, invoked in interacting with a database, using JDBC. Handling exceptions and closing the connection in the right way and in the right sequence has tended to be error-prone. Using `try-with-resources`, it is straightforward to write code in the correct way.

See Demo `lesson10.lecture.jdbc.read_trywithres`. To run the code, you need to set up JavaDB. To do this, see `setup/setting-up-javadb`



# Resources in Java 8 Which Can Be Used with try-with-resources

Any object that implements `java.lang.AutoCloseable`, which includes all objects that implement `java.io.Closeable`, can be used as a resource. The classes in Java 8 that implement `AutoCloseable` are listed here:

```
AbstractInterruptibleChannel, AbstractSelectableChannel, AbstractSelector,
AsynchronousFileChannel, AsynchronousServerSocketChannel, AsynchronousSocketChannel,
AudioInputStream, BufferedInputStream, BufferedOutputStream, BufferedReader, BufferedWriter,
ByteArrayInputStream, ByteArrayOutputStream, CharArrayReader, CharArrayWriter,
CheckedInputStream, CheckedOutputStream, CipherInputStream, CipherOutputStream, DatagramChannel,
DatagramSocket, DataInputStream, DataOutputStream, DeflaterInputStream, DeflaterOutputStream,
DigestInputStream, DigestOutputStream, FileCacheImageInputStream, FileCacheImageOutputStream,
FileChannel, FileImageInputStream, FileImageOutputStream, FileInputStream, FileLock,
FileOutputStream, FileReader, FileSystem, FileWriter, FilterInputStream, FilterOutputStream,
FilterReader, FilterWriter, Formatter, ForwardingJavaFileManager, GZIPInputStream,
GZIPOutputStream, ImageInputStreamImpl, ImageOutputStreamImpl, InflaterInputStream,
InflaterOutputStream, InputStream, InputStream, InputStream, InputStreamReader, JarFile,
JarInputStream, JarOutputStream, LineNumberInputStream, LineNumberReader, LogStream,
MemoryCacheImageInputStream, MemoryCacheImageOutputStream, MLet, MulticastSocket,
ObjectInputStream, ObjectOutputStream, OutputStream, OutputStream, OutputStream,
OutputStreamWriter, Pipe.SinkChannel, Pipe.SourceChannel, PipedInputStream, PipedOutputStream,
PipedReader, PipedWriter, PrintStream, PrintWriter, PrivateMLet, ProgressMonitorInputStream,
PushbackInputStream, PushbackReader, RandomAccessFile, Reader, RMIConnectionImpl,
RMIConnectionImpl_Stub, RMIConnector, RMIIIOpsServerImpl, RMIIJRMPServerImpl, RMIServerImpl,
Scanner, SelectableChannel, Selector, SequenceInputStream, ServerSocket, ServerSocketChannel,
Socket, SocketChannel, SSLServerSocket, SSLSocket, StringBufferInputStream, StringReader,
StringWriter, URLClassLoader, Writer, XMLDecoder, XMLEncoder, ZipFile, ZipInputStream,
ZipOutputStream
```

# Main Point 1

Java's new try-with-resources construct handles the complex logic behind exception-handling when the process requires that a resource be "closed" after its use. Earlier versions of Java required developers to handle this situation with an error-prone imperative programming style; Java's new approach is declarative: The developer just tells the runtime which resources will be used and the runtime takes care of the close operations, including exception-handling. This effortless approach to handling code has a parallel in the domain of human life: When individual awareness is sufficiently attuned to this deeper levels of consciousness, very little effort is required to achieve significant success.



# Overview

1. Unit-testing Stream Pipelines
2. Introduction to Annotations
3. Using Java 8's try-with-resources
4. **Handling Exceptions Arising in Stream Pipelines**
5. Concurrent Processing and Parallel Streams

# Handling Exceptions Arising in Stream Pipelines

1. Ordinary functional expressions, composed in a pipeline, may throw exceptions, but very often exception-handling can be done in the usual way. See demo code in `lesson10.lecture.exceptions`.
2. However, stream operations, like `map` and `filter`, that require a functional interface whose unique method *does not have a throws clause* (like `Function` and `Predicate`), make exception-handling more difficult. See demo code to see issues and best possible solutions.  
`lesson10.lecture.exceptions2`
3. The best one can do in these situations is to convert checked exceptions to `RuntimeExceptions`. The code can be made more readable and compact if the `try/catch` clause that is needed can be tucked away in an auxiliary method. Examples are provided in `lesson10.lecture.exceptions3`, `lesson10.lecture.exceptions.connectold`, and `lesson10.lecture.exceptions.connectnew`



# Overview

1. Unit-testing Stream Pipelines
2. Introduction to Annotations
3. Using Java 8's try-with-resources
4. Handling Exceptions Arising in Stream Pipelines
5. **Concurrent Processing and Parallel Streams**

# Concurrent Processing and Parallel Streams - Overview

- A. Introduction to threads
- B. Working with threads: the Runnable interface
- C. Thread safety and the synchronized keyword
- D. Java 8 convenience class for invoking threads: the Executor class
- E. When should you use parallel streams?



# Introduction to Threads

1. A *process* in Java is an instance of a Java program that is being executed. It contains the program code and its current activity. A process has a self-contained execution environment. A process generally has a complete, private set of basic run-time resources; in particular, each process has its own memory space.
2. A *thread* is a component of a process. Multiple threads can exist within the same process, executing concurrently (one starting before others finish) and sharing memory (and other resources), while different processes do not share these resources. In particular, the threads of a process share the values of its variables at any given moment.
3. Every process has at least one thread, the *main thread* (the main method of a Java program starts up the main thread.) Other threads may be created from the main thread.

# Introduction to Threads (cont.)

4. Multiple threads are typically invoked to perform multiple tasks simultaneously, or to simulate simultaneous execution of multiple tasks. In a multiprocessor environment, different threads can access different processors; in a single processor environment, multiple threads can appear to work simultaneously by virtue of *time-slicing* – the operating system allots portions of time to competing threads.
5. Examples of how multiple threads are used:
  - a. One thread keeps a UI active while another thread performs a computation or accesses a database
  - b. Divide up a long computation into pieces and let each thread compute values for one piece, then combine the results (computing in parallel)
  - c. Web servers typically handle client requests on separate threads; in this way, many clients can be served “simultaneously.”



# Creating Threads in Java

Code that you wish to run in a new thread is contained in the `run()` method of a class that implements the `Runnable` interface.

```
interface Runnable {  
    void run() ;  
}  
class MyRunnable implements Runnable {  
    public void run() {  
        System.out.println("Running a thread!");  
    }  
}
```

# Creating Threads in Java (cont.)

The thread is then *spawned* when an instance of your class is used as an argument to the Thread constructor, and the start() method is called on the Thread instance.

The following code creates a thread and starts it:

```
MyRunnable myRunnable = new MyRunnable();  
Thread t = new Thread(myRunnable);  
t.start();
```



# Testing Singleton Using a Single Thread

```
public class Singleton {  
    private static Singleton instance;  
    public static int counter = 0;  
    private Singleton() {  
        incrementCounter();  
    }  
    public static Singleton getInstance() {  
        if(instance == null) {  
            instance = new Singleton();  
        }  
        return instance;  
    }  
    private static void incrementCounter() {  
        counter++;  
    }  
}
```

```
public class SingleThreadedTest2 {  
    public static void main(String[] args) {  
        for(int i = 0; i < 1; ++i) {  
            createAndStartThread();  
            System.out.println("Num instances: " + Singleton.counter);  
        }  
    }  
    public static void createAndStartThread() {  
        Runnable r = () -> {  
            for(int i = 0; i < 1000; ++i) {  
                Singleton.getInstance();  
            }  
        };  
        new Thread(r).start();  
        try {  
            Thread.sleep(10);  
        } catch (InterruptedException e) {}  
    }  
}
```

As expected, only 1 instance is ever created.

**Note:** We have put each thread to sleep for 10 milliseconds before allowing the next one to start. If we do not do this, then the first 1 or 2 calls of `createAndStart` will record *0 instances*. This is because the change made by each thread may not be visible to the main thread immediately (this is most likely because processor memory is much faster than the RAM where the counter data is stored).

# Testing Singleton Using Multiple Threads

```
public class Singleton {  
    private static Singleton instance;  
    public static int counter = 0;  
    private Singleton() {  
        incrementCounter();  
    }  
    public static Singleton getInstance() {  
        if(instance == null) {  
            instance = new Singleton();  
        }  
        return instance;  
    }  
    private static void incrementCounter() {  
        counter++;  
    }  
}
```

```
public class MultiThreadedTest {  
    public static void main(String[] args) {  
        for(int i = 0; i < 10; ++i) {  
            multipleCalls();  
            System.out.println("Num instances: " + Singleton.counter);  
        }  
    }  
    public static void multipleCalls() {  
        Runnable r = () -> {  
            for(int i = 0; i < 5000; ++i) {  
                Singleton.getInstance();  
            }  
        };  
        for(int i = 0; i < 1000; ++i) {  
            new Thread(r).start();  
        }  
    }  
}
```

The test shows that competing threads are creating multiple instances of the Singleton class. The test “instance == null” is being interrupted so that it appears to be true to more than one thread, and so the constructor is called multiple times.



# Race Conditions and Thread Safety

1. When two or more threads have access to the same object and each modifies the state of the object, this situation is called a *race condition*, which arises when threads *interfere* with each other (the sequence of steps being executed by one thread is interrupted by another thread).
2. Code is said to be *thread-safe* if it is guaranteed to be free of race conditions when accessed by multiple threads simultaneously.
3. We can say therefore that this Singleton implementation is *not thread-safe*.

# Forcing Serialized Access with synchronized

1. We can force threads to access the getInstance method of Singleton *one at a time* (this is called *serialized access*) by labeling getInstance with the keyword synchronized.
2. When a method is synchronized, in order for a thread to execute the method, it must *acquire the lock* for the instance of the object that the method is running on. Each object has an *intrinsic* lock, and a thread gains access to this lock when it calls the method, as long as no other thread has the lock. Once a thread executes the synchronized method, the lock becomes available again and the next eligible thread (determined by the OS using thread priorities and other factors) then acquires the lock.



# Forcing Serialized Access with synchronized (cont.)

**Note:** This use of the word “serialized” has nothing to do with the Serializable interface that we examined earlier in the course

**Note:** We have seen that competing threads could corrupt the “== null” test. However, competing threads could also corrupt the counter since the increment operation `counter++` is not atomic (it is in fact the assignment `counter = counter + 1`). Therefore, to be sure that the `MultiThreadedTest` is really producing multiple instances of `Singleton`, we must make the `incrementCounter` method synchronized, as in the code below. Running this test, and witnessing multiple calls to the `Singleton` constructor once again convinces us that multiple instances are being created.

```
public class Singleton2 {
    private static Singleton2 instance;
    public static int counter = 0;
    private Singleton2() {
        incrementCounter();
    }
    public static Singleton2 getInstance() {
        if(instance == null) {
            instance = new Singleton2();
        }
        return instance;
    }
    /* Guarantees proper count of instances */
    synchronized private static void incrementCounter() {
        counter++;
    }
}
```

# Exercise 10.4

The code to the right is an implementation of the Queue data type. Are its operations threadsafe? If not, think of a test you could do, involving multiple threads, that would prove that it is not. (You do not need to write any code, but you will in the lab, so you will be able to use the ideas you come up with now.)

```
public class Queue {  
    class Node {  
        Object value;  
        Node next;  
    }  
    private Node head;  
    private Node tail;  
    public void add(Object newValue) {  
        Node n = new Node();  
        if(head == null) head = n;  
        else tail.next = n;  
        tail = n;  
        tail.value = newValue;  
    }  
    public Object remove() {  
        if(head == null) return null;  
        Node n = head;  
        head = n.next;  
        return n.value;  
    }  
}
```



# Main Point 2

Code that will perform flawlessly in a single-threaded environment may bring disastrous consequences in a multi-threaded environment if shared data is not managed properly.

In life, a similar situation exists because the particulars of life are always changing. In any domain of life – work, family, health, society – a situation that appears favorable today may become unfavorable tomorrow. How does one prepare for change that cannot be predicted? The only way is to be established in a part of oneself that does not change, no matter how the world changes. That part is first experienced as simple silence in the practice of TM. In time, with regular practice, it becomes firmly established.

# Starting and Managing Threads with Executor

1. When many threads need to be managed, starting threads in the way shown above is not optimal. To create and manage threads properly, Java has an `Executor` class.
2. Two examples of specialized `Executor` classes:
  - `Executors.newCachedThreadPool()` - optimized for creation of threads for performing many small tasks or for tasks which involve long wait periods.
  - `Executors.newFixedThreadPool(numThreads)` – for computationally intensive tasks



# Starting and Managing Threads with Executor (cont.)

3. We modify our earlier code to make use of this the Executor class. We synchronize the getInstance method in SynchronizedSingleton.

```
public class SynchronizedSingleton {  
    private static SynchronizedSingleton instance;  
    public static int counter = 0;  
    private SynchronizedSingleton() {  
        incrementCounter();  
    }  
    synchronized public static SynchronizedSingleton getInstance() {  
        if(instance == null) {  
            instance = new SynchronizedSingleton();  
        }  
        return instance;  
    }  
    private static void incrementCounter() {  
        counter++;  
    }  
}
```

```
public class MultiThreadedTestWithExec {  
    private static Executor exec  
        = Executors.newCachedThreadPool();  
    public static void main(String[] args) {  
        for(int i = 0; i < 10; ++i) {  
            multipleCalls();  
            System.out.println("Num instances: "  
                + SynchronizedSingleton.counter);  
        }  
    }  
    public static void multipleCalls() {  
        Runnable r = () -> {  
            for(int i = 0; i < 500; ++i) {  
                SynchronizedSingleton.getInstance();  
            }  
        };  
        for(int i = 0; i < 100; ++i) {  
            exec.execute(r);  
        }  
        try {  
            Thread.sleep(10);  
        } catch (InterruptedException e) {}  
    }  
}
```

Note: You may notice that the program waits a bit after the last printout. It terminates when the pooled threads have been idle for a while; after some time, the executor terminates them.

# Guidelines for Using Parallel Streams

1. When you create a parallel stream from a Collection class in order to process elements in parallel, Java handles this request by partitioning the collection and processing each piece with a separate thread.

Not every Stream operation, nor every underlying collection type, works well with parallel processing. We give general guidelines for choosing between sequential and parallel streams.



# Guidelines for Using Parallel Streams (cont.)

## 2. Some Guidelines

- A. Don't use parallel streams if the number of elements is small – the improved performance (if any) will not in this case outweigh the overhead cost of working with parallel streams.
- B. Operations that depend on the order of elements in the underlying collection should not be done in parallel. Example: `findFirst`, `limit`.
- C. If the terminal operation of the stream is expensive (example: `collect(Collectors.joining)`), you must remember that it will be executed repeatedly in a parallel computation – this could be a good reason to avoid parallel streams in this case.
- D. Translation between primitives and their object wrappers becomes very expensive when done in parallel. If you are working with primitives, use the primitive streams, like `IntStream` and `DoubleStream`.
- E. Certain data structures can be divided up more efficiently than others – `ArrayList`, `HashSet` and `TreeSet` can be partitioned efficiently, but `LinkedLists` cannot. (In a `LinkedList`, how can a section of the list be extracted and processed separately?)
- F. Until you develop a degree of expertise in working with parallel computations, it is a good idea to benchmark the performance of a pipeline executed in parallel and compare with the performance of the sequential version.

# Sample Benchmarks for Sequential vs Parallel Processing

Warburton, *Java 8 Lambdas* (p. 84) gives an example of a benchmark test that makes a convincing case for choosing parallel processing over sequential processing for a particular task.

## Sequential Version of Code

```
public int serialArraySum() {  
    return albums.stream()  
        .flatMap(Album::getTracks)  
        .mapToInt(Track::getLength)  
        .sum();  
}
```

```
//counts the total number of  
//tracks on all albums
```



# Sample Benchmarks for Sequential vs Parallel Processing (cont.)

## Parallel Version of Code

```
public int parallelArraySum() {  
    return albums.parallelStream()  
        .flatMap(Album::getTracks)  
        .mapToInt(Track::getLength)  
        .sum();  
}
```

Warburton reports that, when running on a quad core Windows machine, when the number of albums was just 10, the sequential version was 8x faster; when number of albums was 100, the two versions were equally fast; when the number of albums was 10,000, the parallel version was 2.5x faster.

# Connecting the Parts of Knowledge With the Wholeness of Knowledge

## *Annotations*

1. Executing a Java program results in algorithmic, predictable, concrete, testable behavior.
  2. Using annotations, it is possible for a Java program to modify itself and interact with itself.
- 
3. *Transcendental Consciousness* is the field self-referral pure consciousness. At this level, only one field is present, continuously in the state of knowing itself.
  4. *Impulses Within the Transcendental Field*. What appears as manifest existence is the result of fundamental impulses of intelligence within the field of pure consciousness. These impulses are ways that pure consciousness acts on itself, interacts with itself.
  5. *Wholeness Moving Within Itself*. In Unity Consciousness, the diversity of creation is appreciated as the play of fundamental impulses of one's own nature, one's own Self.